


An Optimization Layer for Distributed Matrix Computations

Jonah Brown Cohen, Tselil Schramm, and Ben Weitz

Motivation

- Big data companies like Facebook, Netflix, or Google perform large-scale distributed matrix computations
- Computations experience **trade-offs** in accuracy vs. time or money.



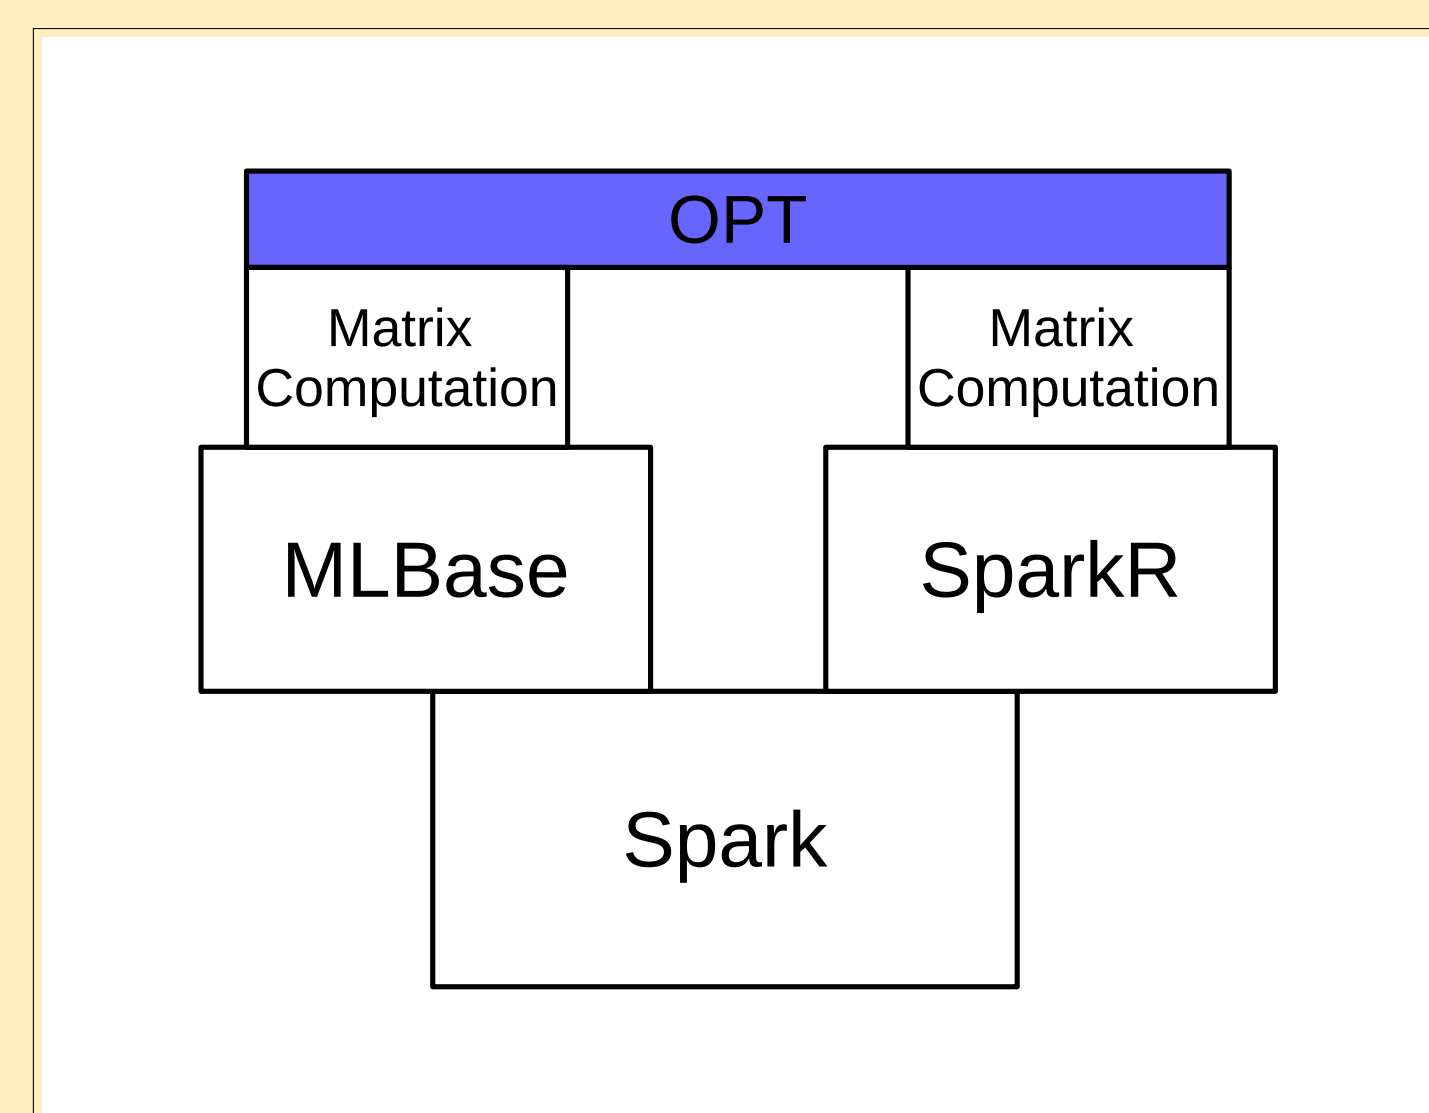
- Human operators manually tweak parameters and partitioning
- Humans are prone to error and costly to hire!
- Solution: Build an optimization layer to automatically tweak and manage these computations
 - Learn to adjust parameters from past computations
 - Incoming jobs come with budgets of time or accuracy that must be met

Objective

- Create an optimizer that automatically picks algorithm parameters and the degree of data partitioning to meet budget specifications

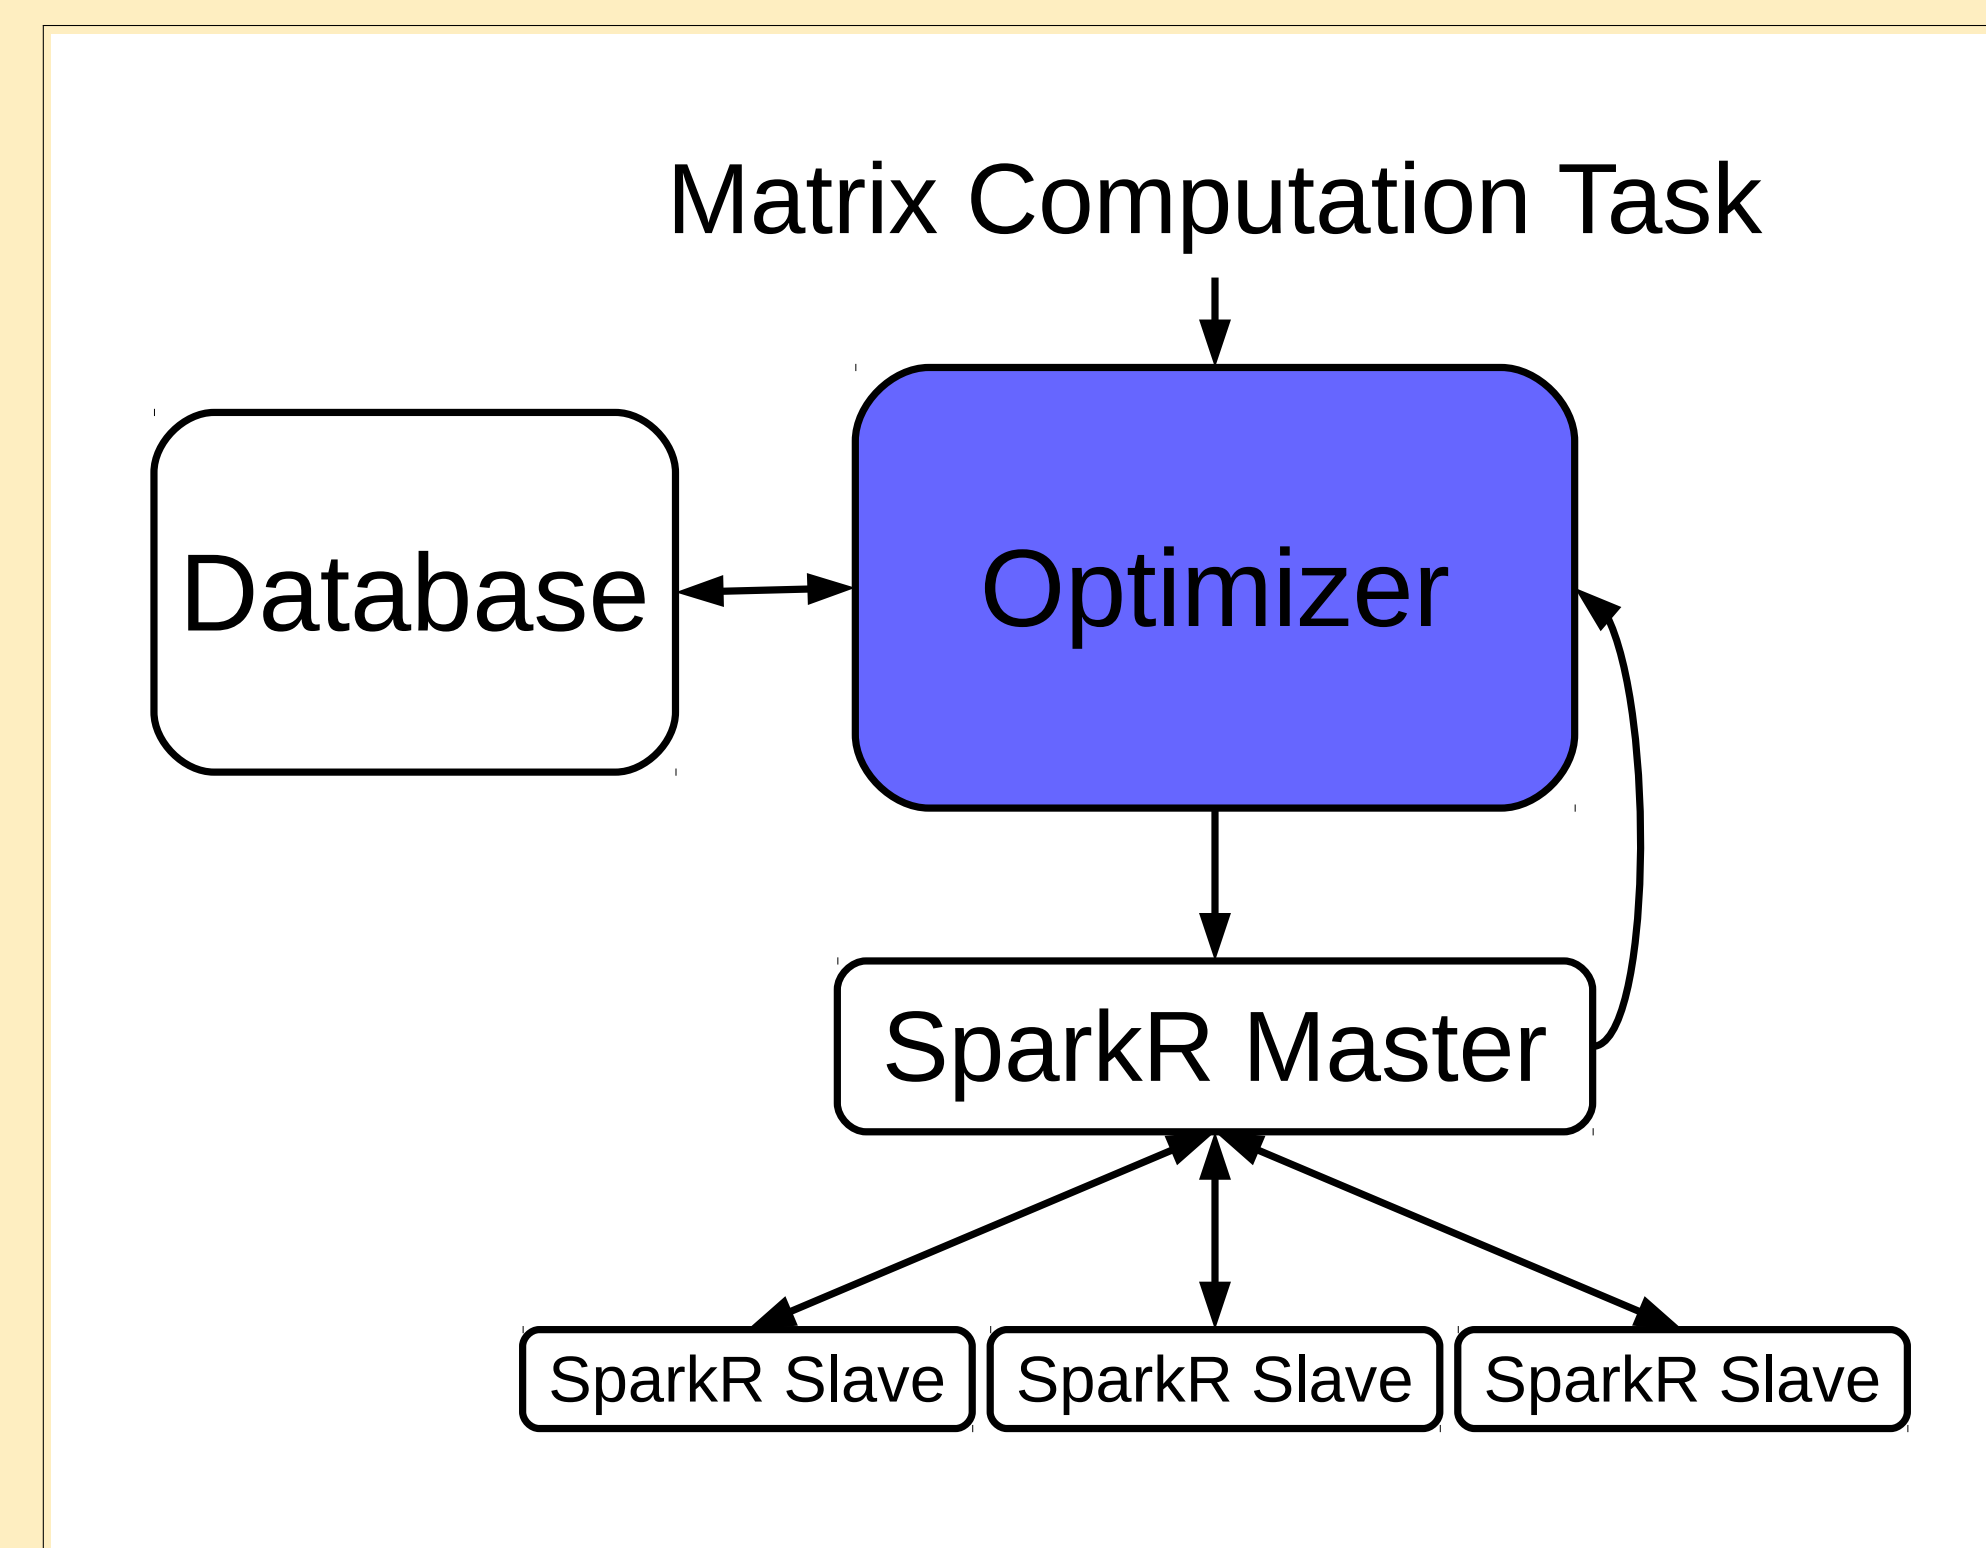
Framework

- Optimizer built on Python
- Interfaces with matrix algorithms implemented in SparkR or MLBase
- Optimizer interfaces directly with algorithms, all parameters hidden from user



Optimizer Design

- Architecture-independent
 - Chooses parameters based on statistics from prior jobs
 - sdf, sdfad, sdf
- Adaptive
 - sdfad, sdf
- Local-optimum Avoiding
 - sdfad, sdf



Implementation

- The words chosen by the adversary are hard,
 - But once you know they're difficult it's easy to adjust.
- Top 12 words (probability of losing in 6 turns):

Evaluation

Tested on Synthetic and Real-world Data

- Gaussian Random Matrices
 - Trained on eight random matrices and tested on a ninth

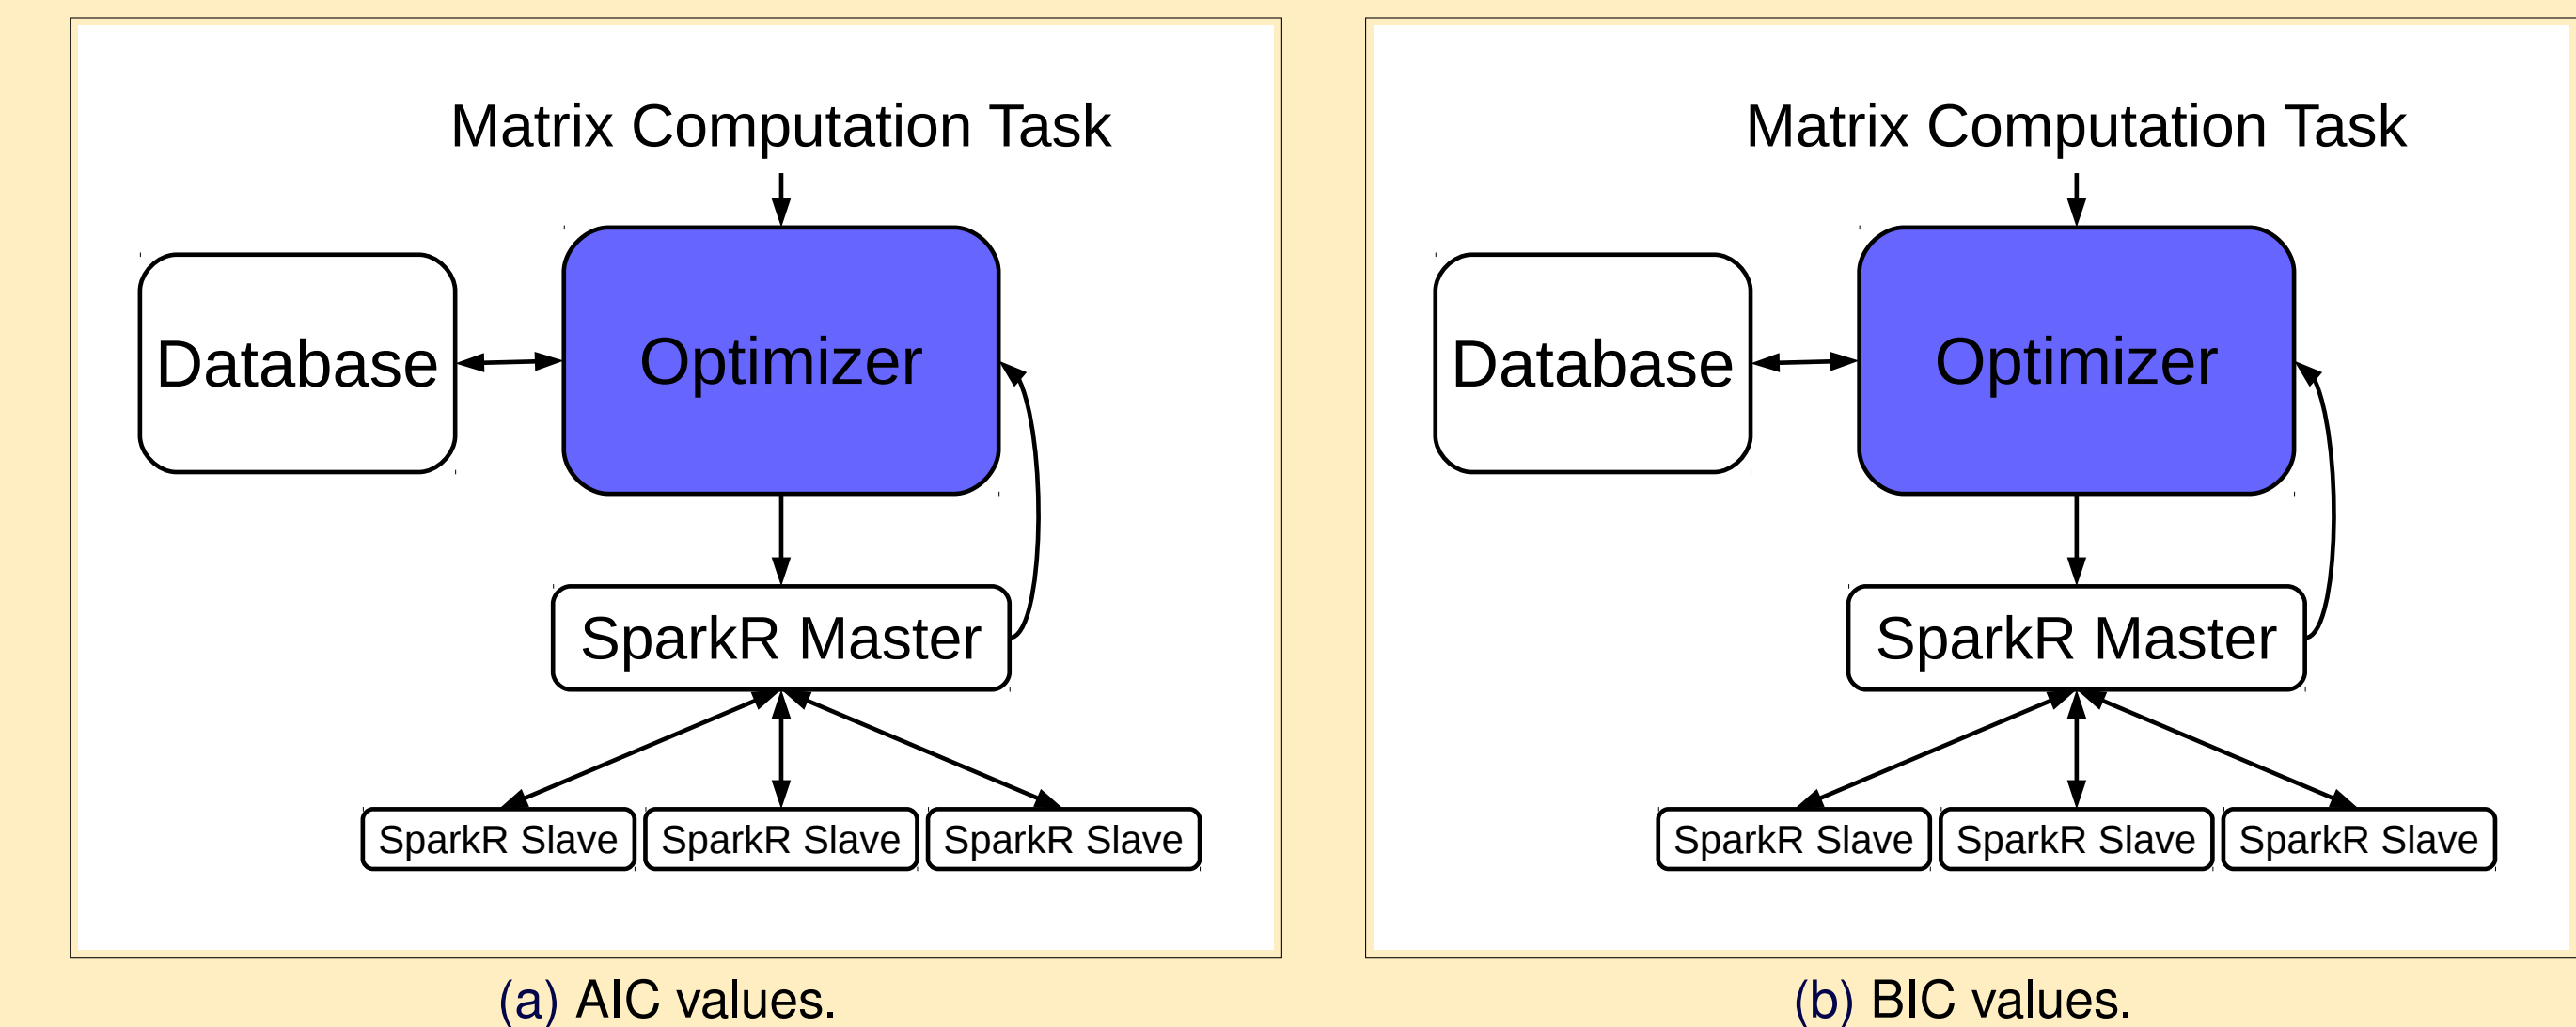


Figure : Plots of AIC and BIC Values against the Memory Parameter for a Dictionary Player with True Memory 2

- MovieLens 10M Dataset
 - Partitioned dataset into 7 parts
 - Trained on all but one partition and tested on the remainder

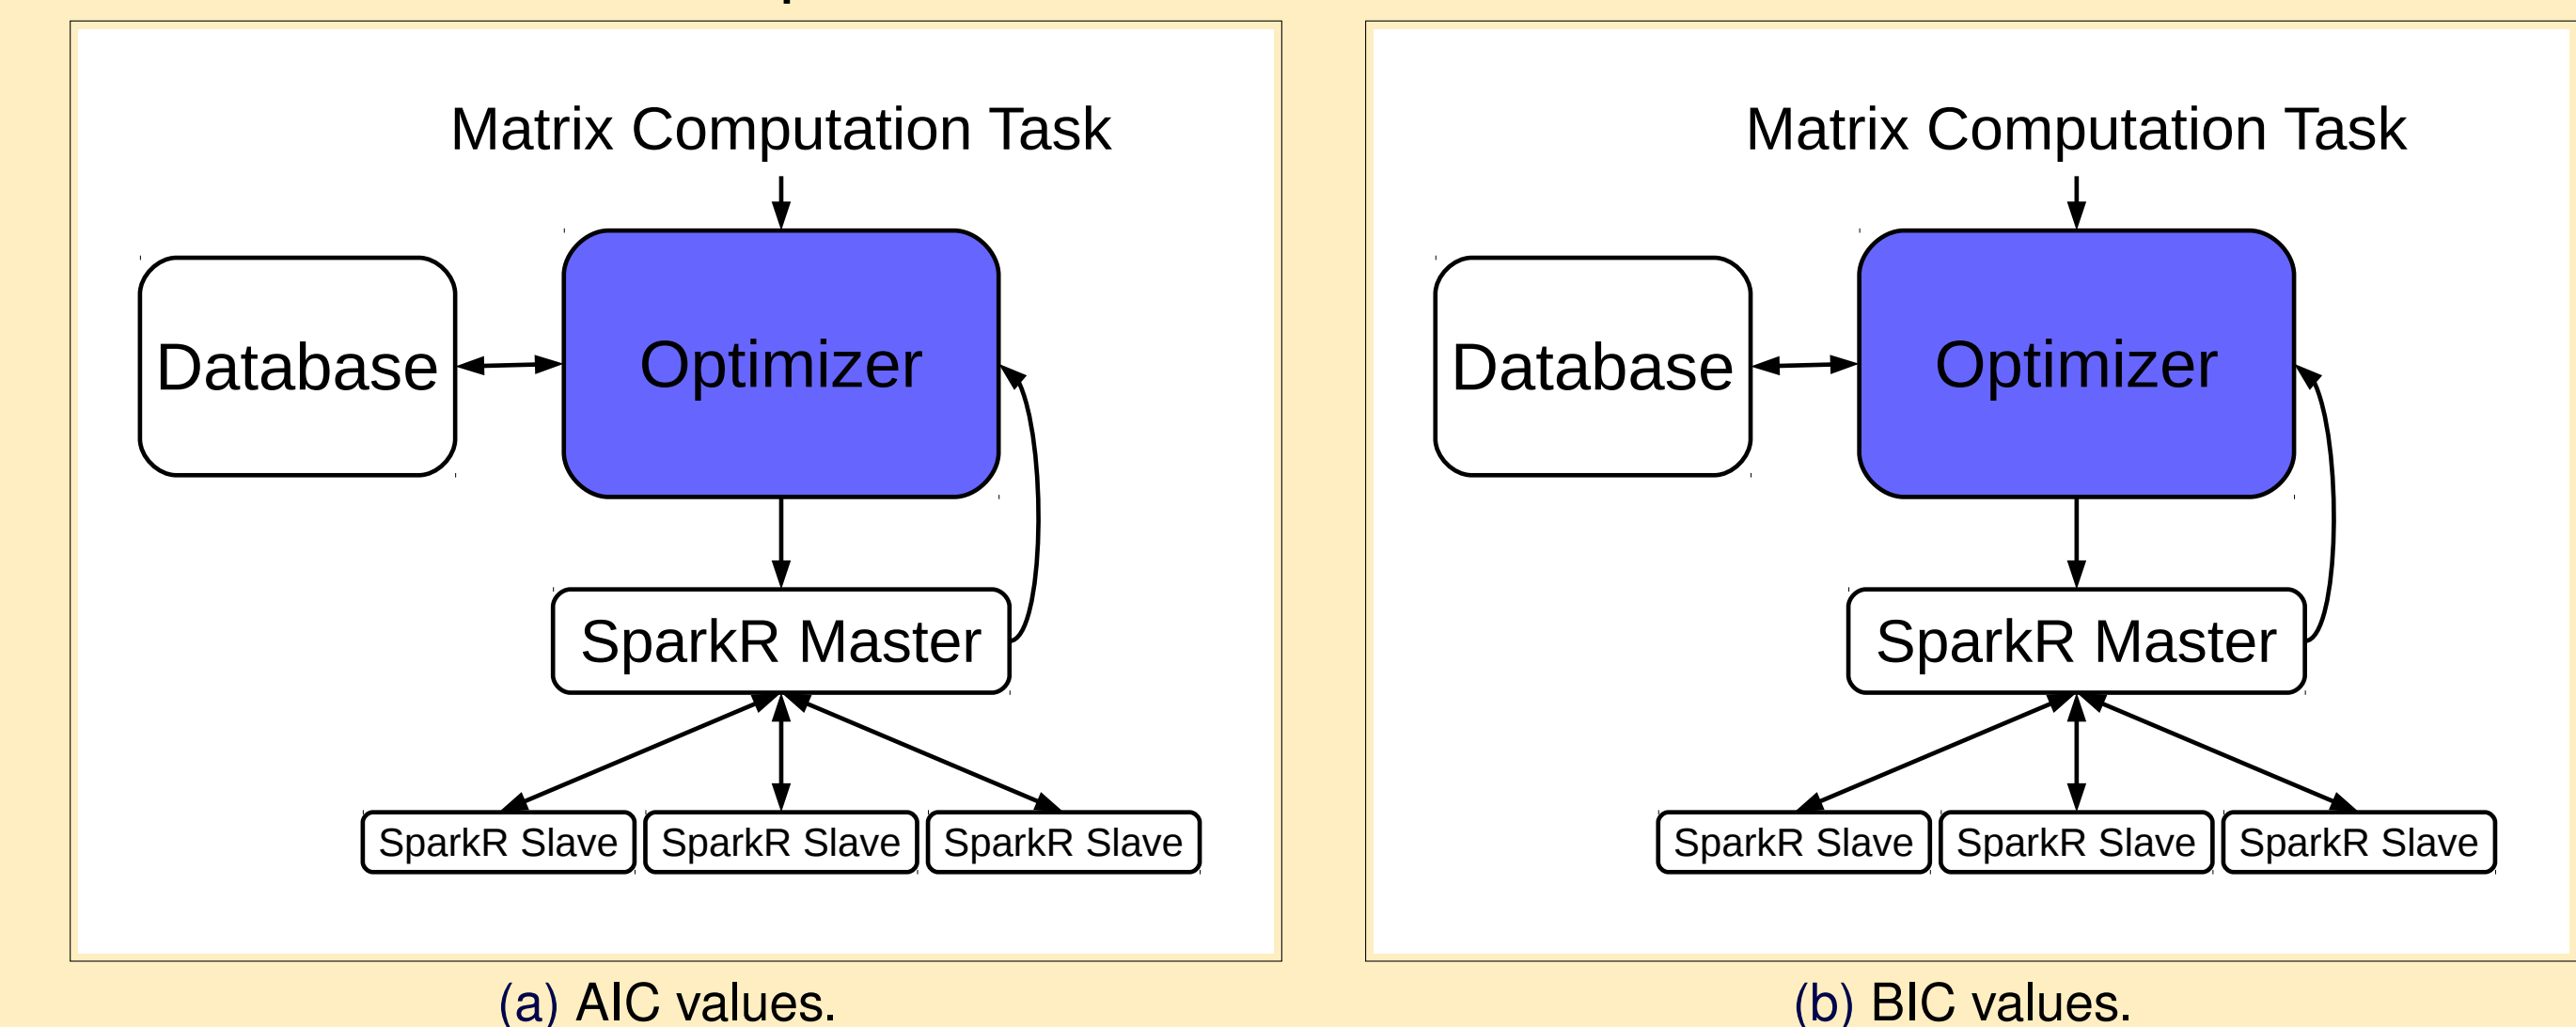


Figure : Plots of AIC and BIC Values against the Memory Parameter for a Dictionary Player with True Memory 2

- Optimizer performs as well as manually setting the parameter!

Future Work

- Optimize over space of algorithms in addition to space of parameters
- Avoid RAM bottlenecks by distributing collect step
- Handle novel or different jobs

Acknowledgements

We would like to thank:

- Shivaram Venkataraman, Ameet Talwalkar, Prof. Anthony Joseph, and Prof. John Kubiawicz
- UC Berkeley and the NSF for providing funds and resources.