

An Optimization Layer for Distributed Matrix Computations

Jonah Brown Cohen, Tselil Schramm, and Ben Weitz

Motivation

- Big data companies like Facebook, Netflix, or Google perform large-scale distributed matrix computations
- Computations experience **trade-offs** in accuracy vs. time or money.




Figure : CAPTION HERE AAARGGHHHH

- Human operators manually tweak parameters and partitioning
- Humans are prone to error and costly to hire!
- Solution: Build an optimization layer to automatically tweak and manage these computations
- Learn to adjust parameters from past computations
- Incoming jobs come with budgets of time or accuracy that must be met

Objective

- Create an optimizer that automatically picks algorithm parameters and the degree of data partitioning to meet budget specifications**

Framework

- Optimizer interface to user: input data, methods to run, and time, error, and financial budgets
- Optimizer then interfaces with matrix algorithms implemented within a parallel framework (such as an algorithm written in SparkR or methods from MLBase)
- All parameter selection hidden from user

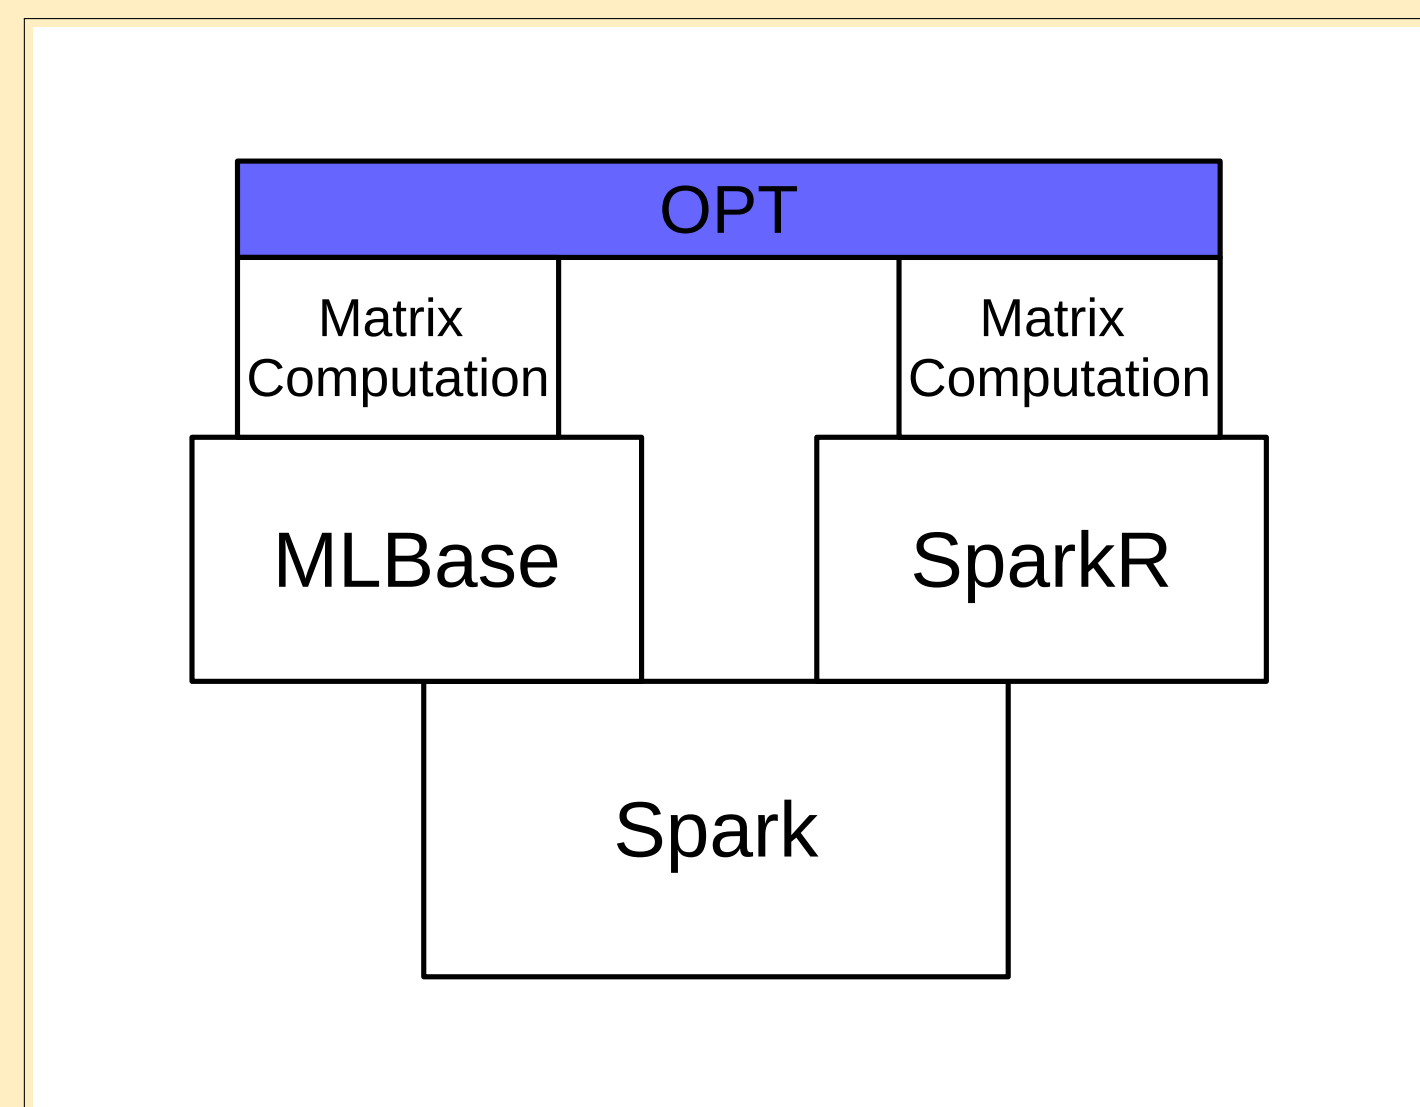


Figure : Our optimizer provides an interface to the user, through which the user specifies the algorithm and a time, error, and monetary budget. The optimizer then calls down to the matrix computation framework with the appropriate parameters.

Optimizer Design

Our optimizer is:

- Architecture-independent
 - Stores statistics from prior jobs on the architecture in question
 - Parameters chosen based on statistics for the specific architecture
- Adaptive
 - Stores statistics from previous jobs on instances with similar distribution
 - Prediction of the optimal parameters improves as the optimizer learns
- Avoids Local Optima
 - When one choice of parameters already encountered is slightly better, there is a risk of getting stuck at a local optimum
 - In “explore mode” we choose the instance parameters randomly, with probability proportional to the relative suitability of the parameters

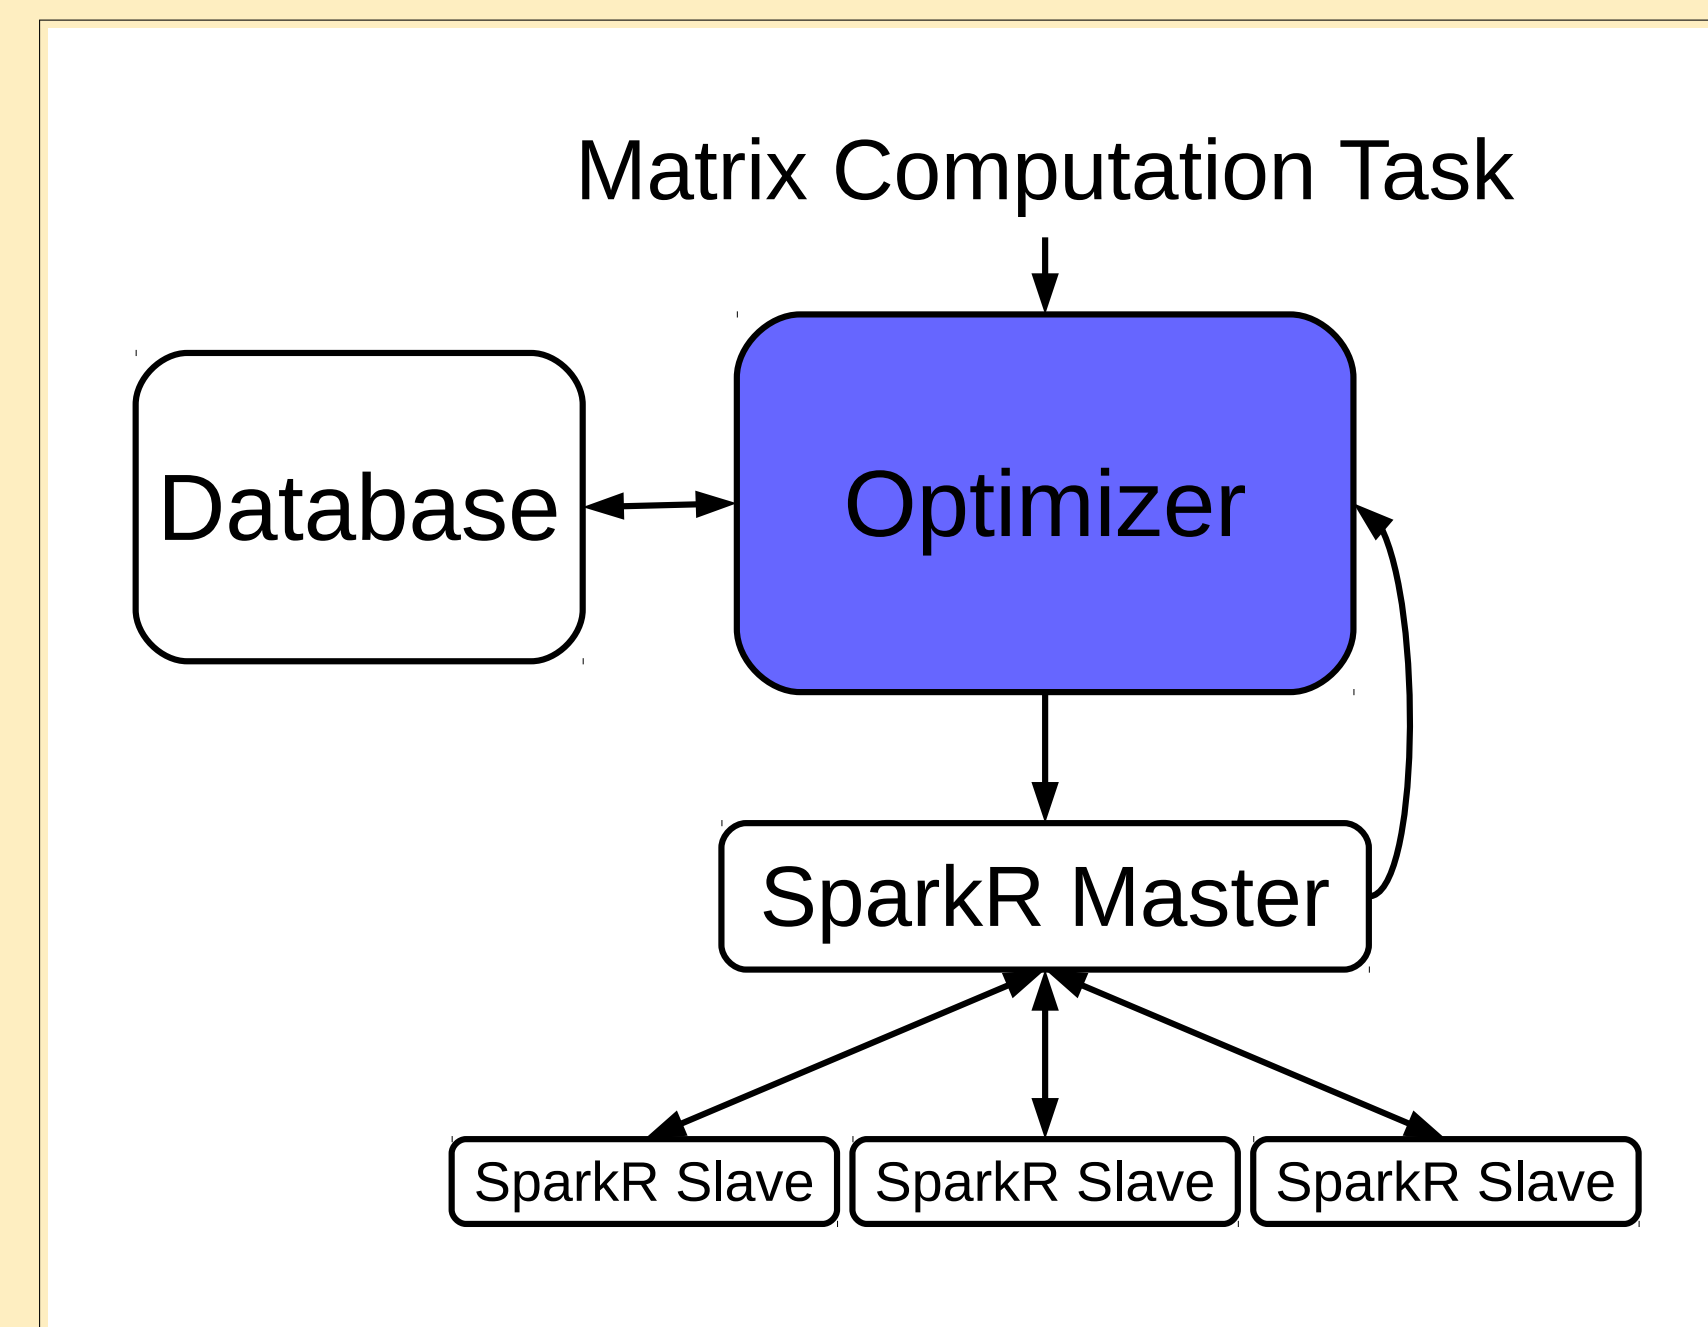


Figure : The control flow of our optimizer. The optimizer looks up relevant information in a database, then interacts with the distributed computation framework (in our current implementation, SparkR).

Implementation

- The matrix computation: Divide-Factor-Combine (DFC)

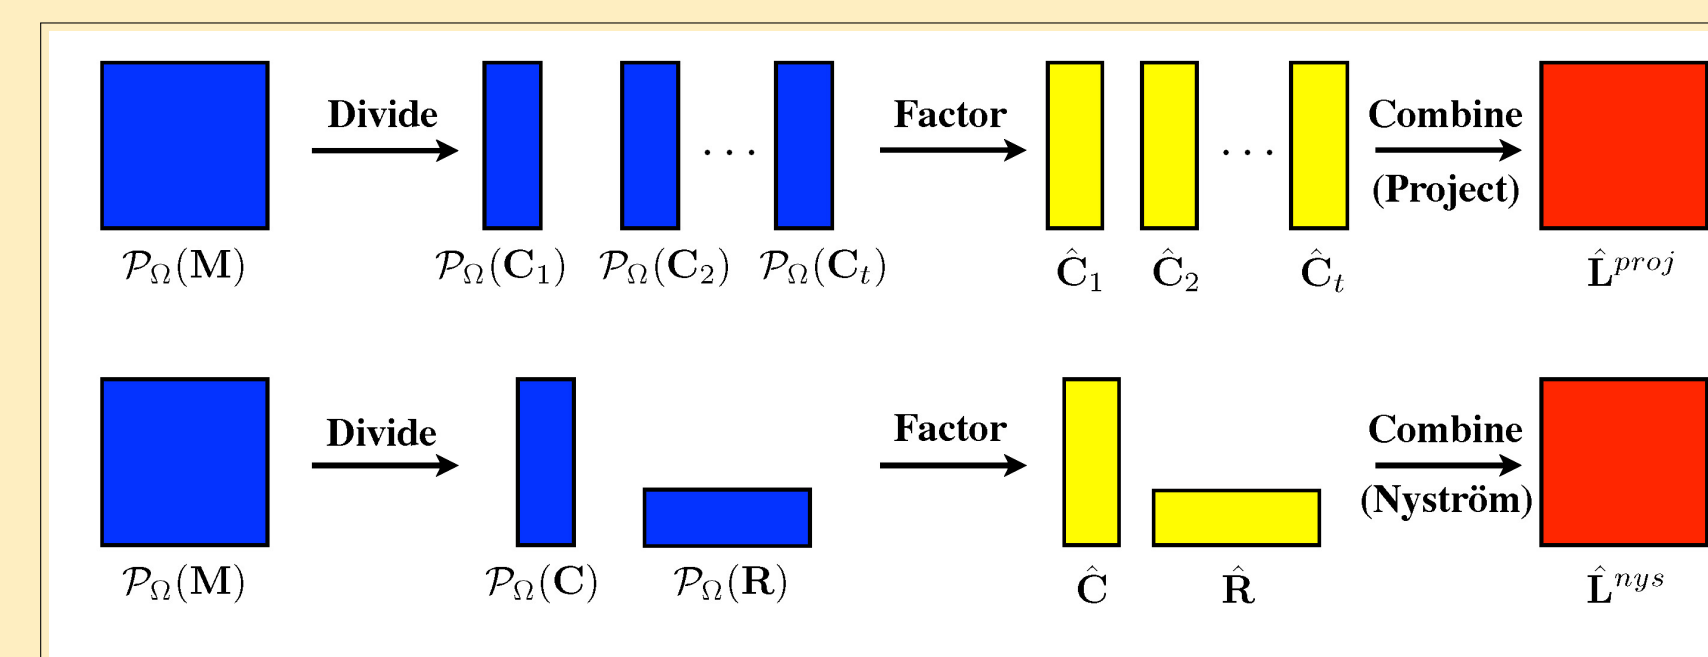


Figure : DFC is a distributed matrix factorization algorithm.

- The parallel computation framework: SparkR
 - An interface for Spark in R
- Implemented DFC to be incorporated into SparkR
- Both randomized and deterministic projection methods
- Two different base factorization methods: APG and SGD


Evaluation

Tested on Synthetic and Real-world Data

- Gaussian Random Matrices
 - Trained on eight random matrices and tested on a ninth



(a) Training Data

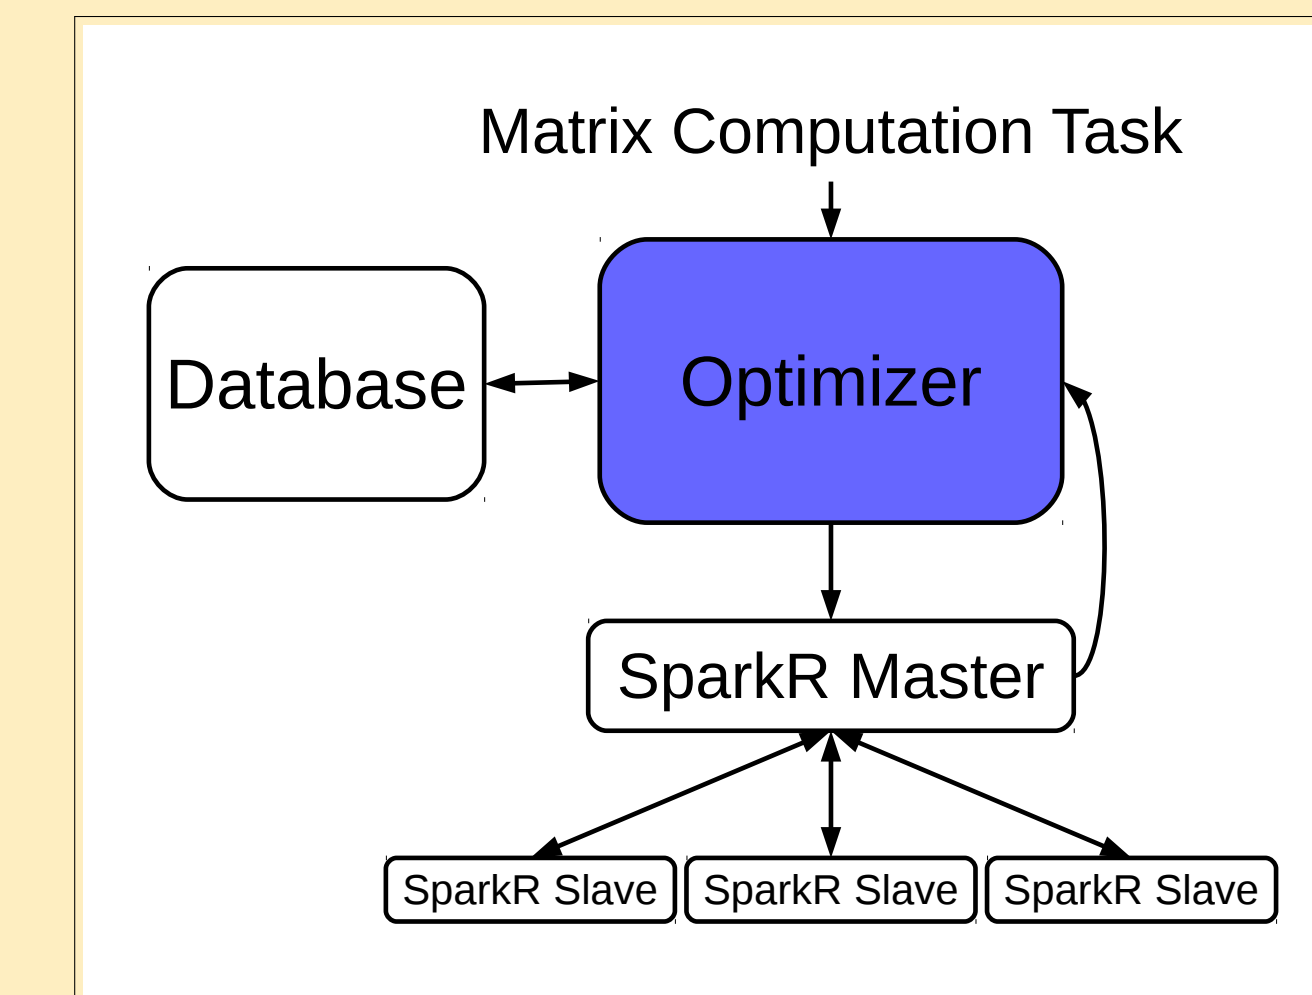


(b) Testing Data

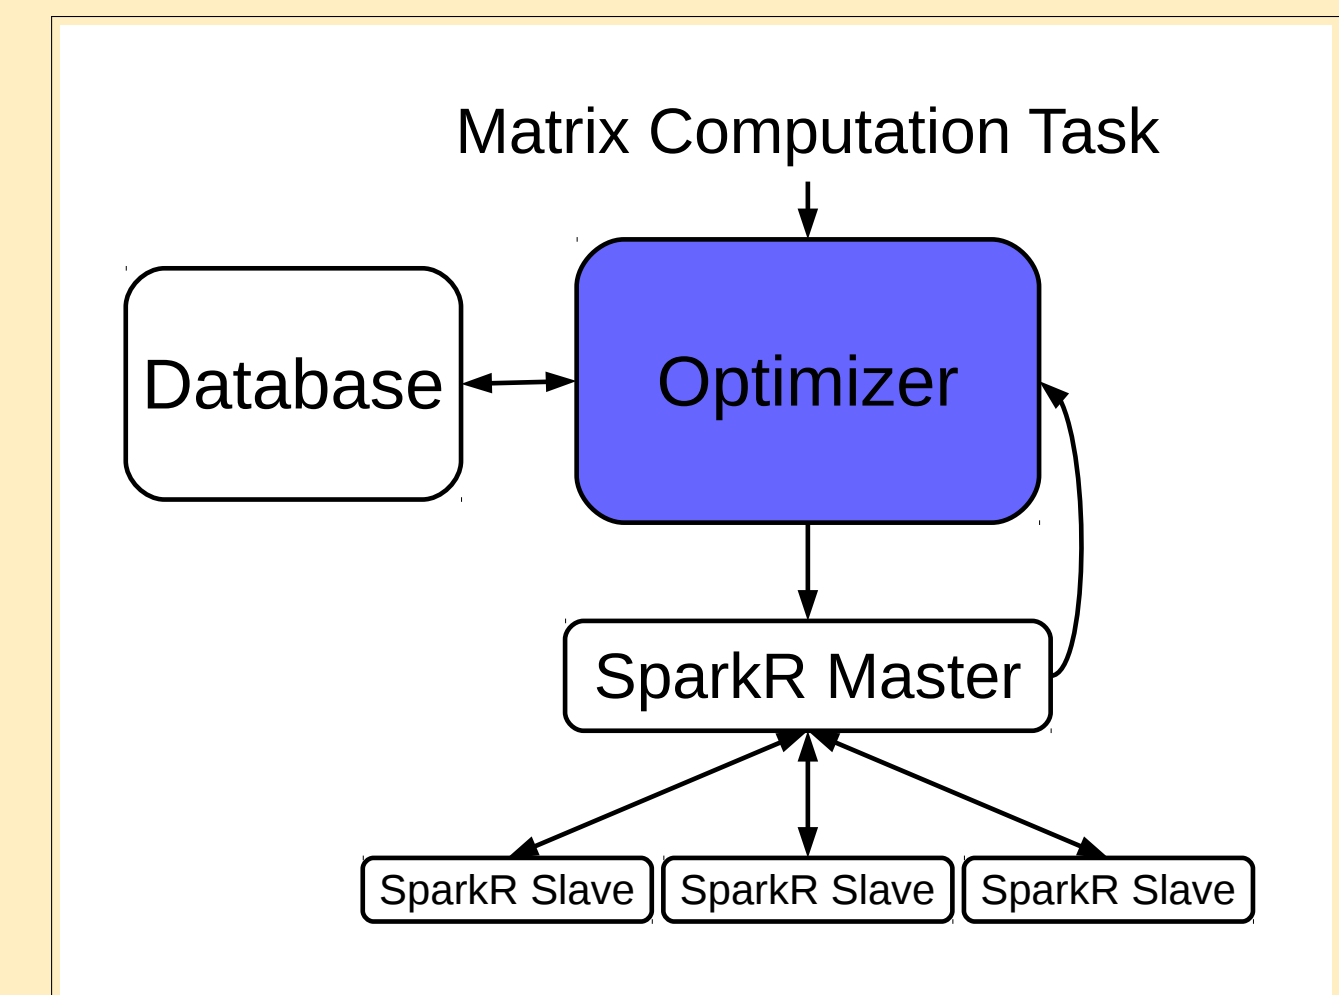
Figure : Plots of Time and Error taken to factor a 4000 by 4000 matrix drawn from a Gaussian distribution

- MovieLens 10M Dataset

- Partitioned dataset into 7 parts
- Trained on all but one partition and tested on the remainder



(a) Training Data



(b) Testing Data

Figure : Plots of Time and Error taken to factor a 4000 by 4000 matrix drawn from a Gaussian distribution

- Optimizer performs as well as manually setting the parameter!

Future Work

- Optimize over space of algorithms in addition to space of parameters
- Avoid RAM bottlenecks by distributing collect step
- Handle novel or different jobs

Acknowledgements

We would like to thank:

- Shivaram Venkataraman, Ameet Talwalkar, Prof. Anthony Joseph, and Prof. John Kubiawicz
- UC Berkeley and the NSF for providing funds and resources.